

---

# Convolutional Neural Networks

Palma Project

---

# Building the CNN

# Steps of the CNN model:

**STEP 1:** Convolution



**STEP 2:** Max Pooling



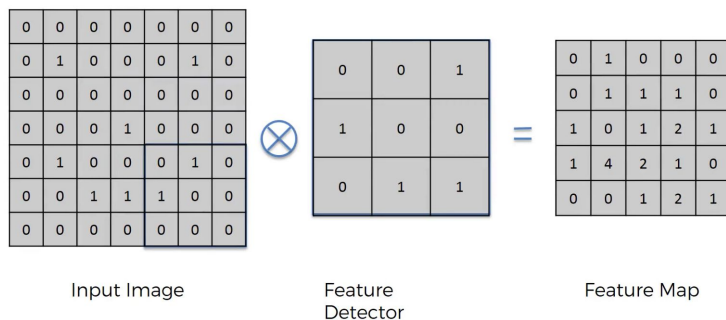
**STEP 3:** Flattening



**STEP 4:** Full Connection

# Step 1 - Convolution

**Purpose**: Find features in the image using a feature detector, put them in a feature map but still preserve the special relationships between pixels.



Reading: Convolution Matrix

# Step 1 - Convolution

**Example:** Feature map of Geoffrey Hinton using a feature detector



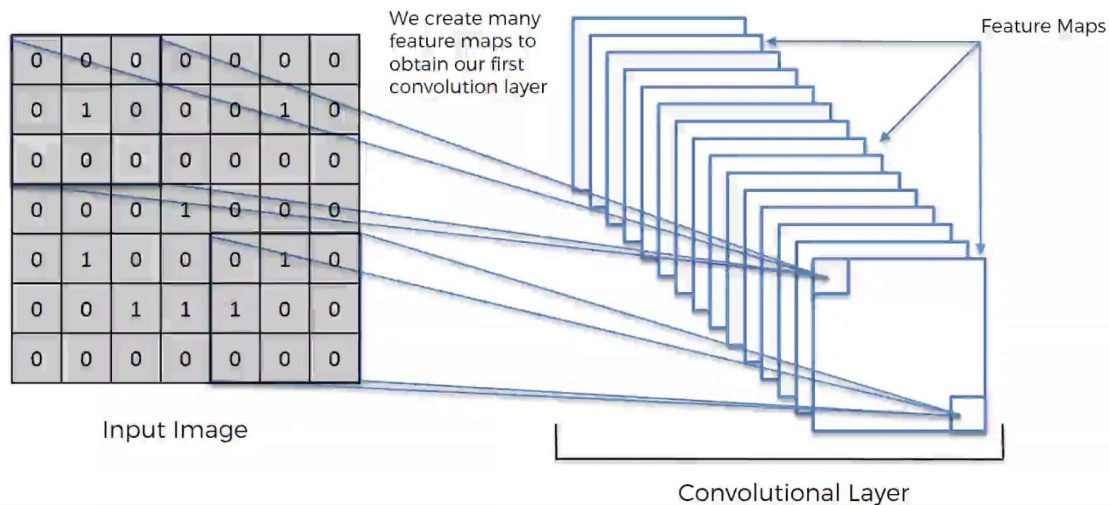
\*

1	0	-1
2	0	-2
1	0	-1



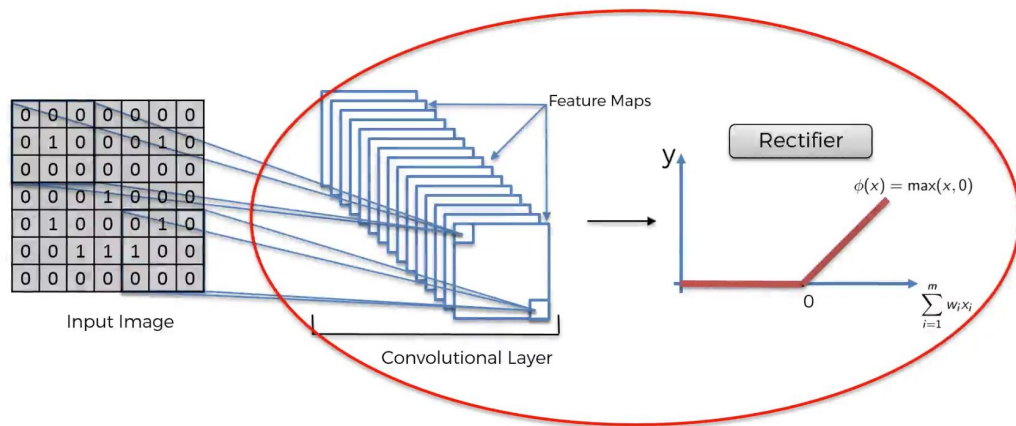
# Step 1 - Convolution

**Purpose:** By using different Feature Detectors, we create many feature maps and then the first convolution layer. Each feature map corresponds to one specific feature of the image



# Step 1 (B) - ReLU Layer

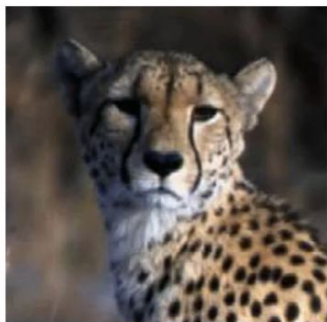
**Purpose:** Increase non linearity because images themselves are highly non linear.



[Reading: Understanding Convolutional Neural Networks with A Mathematical Model](#)

## Step 2 - Pooling

We want the neural network to recognize the image even if it's rotated or a little bit squashed





## Step 2 - Pooling

We use Pooling in order to:

- Introduce Invariance
- Reduce the size of the feature map
- Remove irrelevant information and then prevent **overfitting**

# Step 2 - Pooling

Example of Pooling: **Max Pooling**

0	1	0	0	0
0	1	1	1	0
1	0	1	2	1
1	4	2	1	0
0	0	1	2	1

Feature Map

Max Pooling

1	1	0
4	2	1
0	2	1

Pooled Feature Map

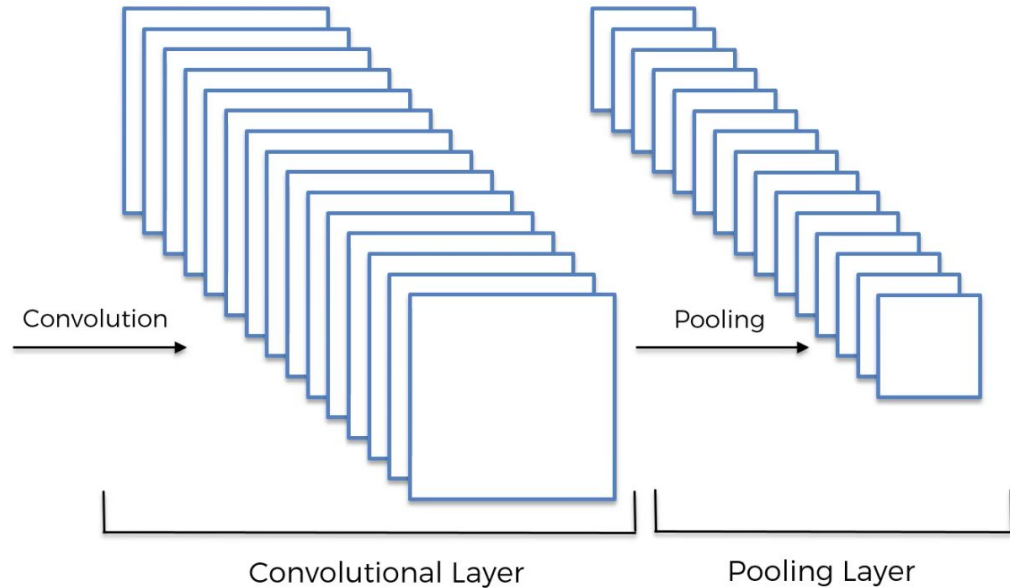
[Reading: Evaluation of Pooling Operations in Convolutional Architectures for Object Recognition](#)

# Step 2 - Pooling

So far..

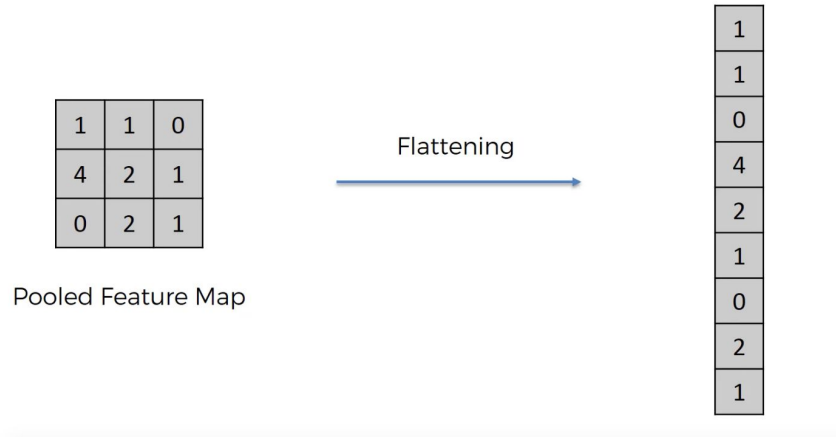
0	0	0	0	0	0	0
0	1	0	0	0	1	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	1	0	0	0	1	0
0	0	1	1	1	0	0
0	0	0	0	0	0	0

Input Image



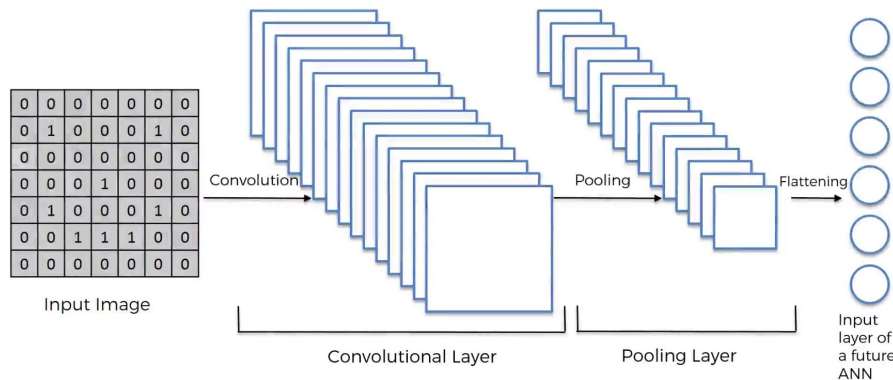
## Step 3 - Flattening

We flatten the Pooled Feature Map by taking the numbers row by row and put them in one column as an input to a future Artificial Neural Network.

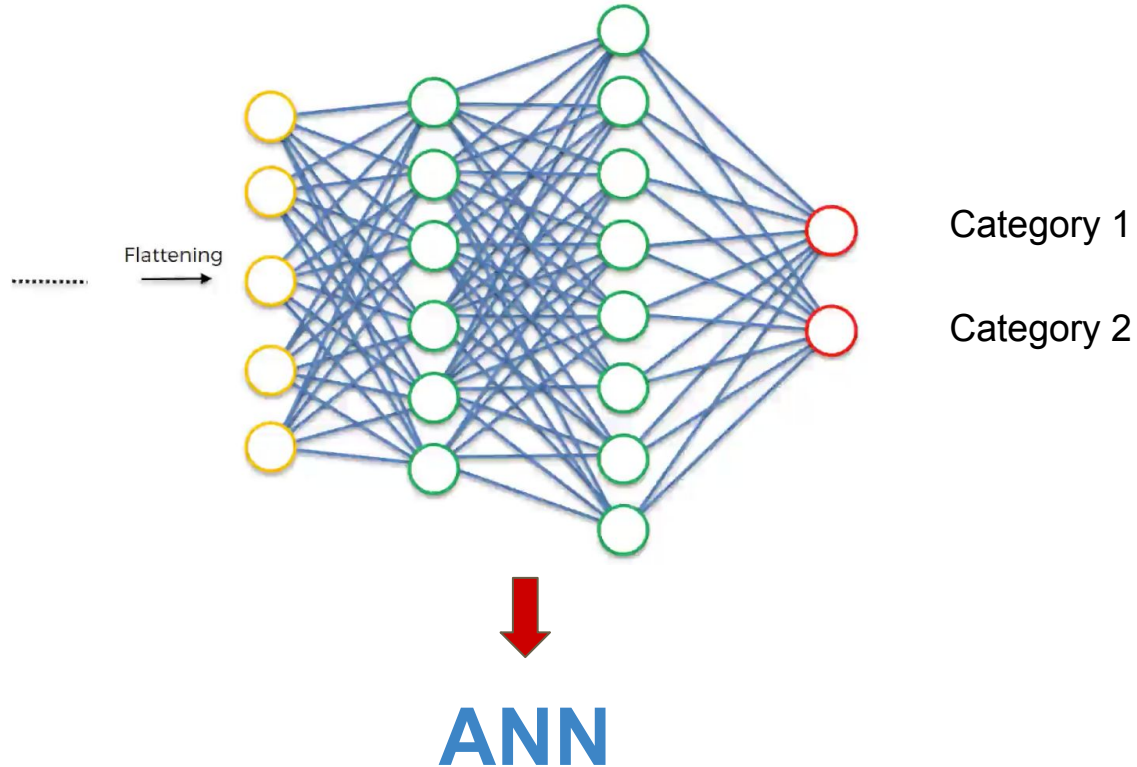


# Step 3 - Flattening

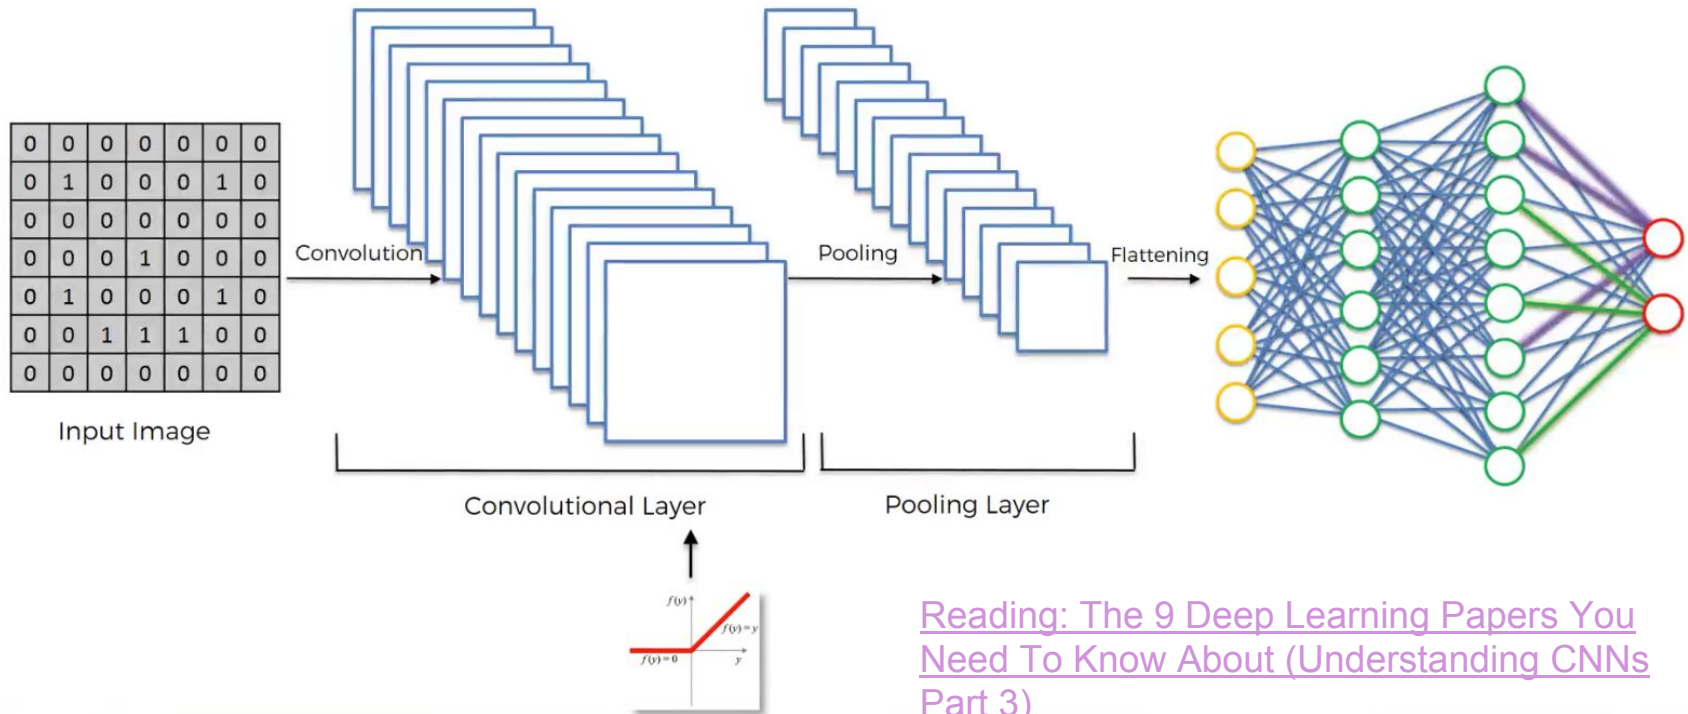
- We flatten all the pooled feature maps into a huge one-dimensional vector.
- Since each feature map corresponds to one specific feature of the image, then each node of this huge vector will represent an information of a specific feature of the input image.



# Step 4 - Full Connection



# Step 4 - Summary

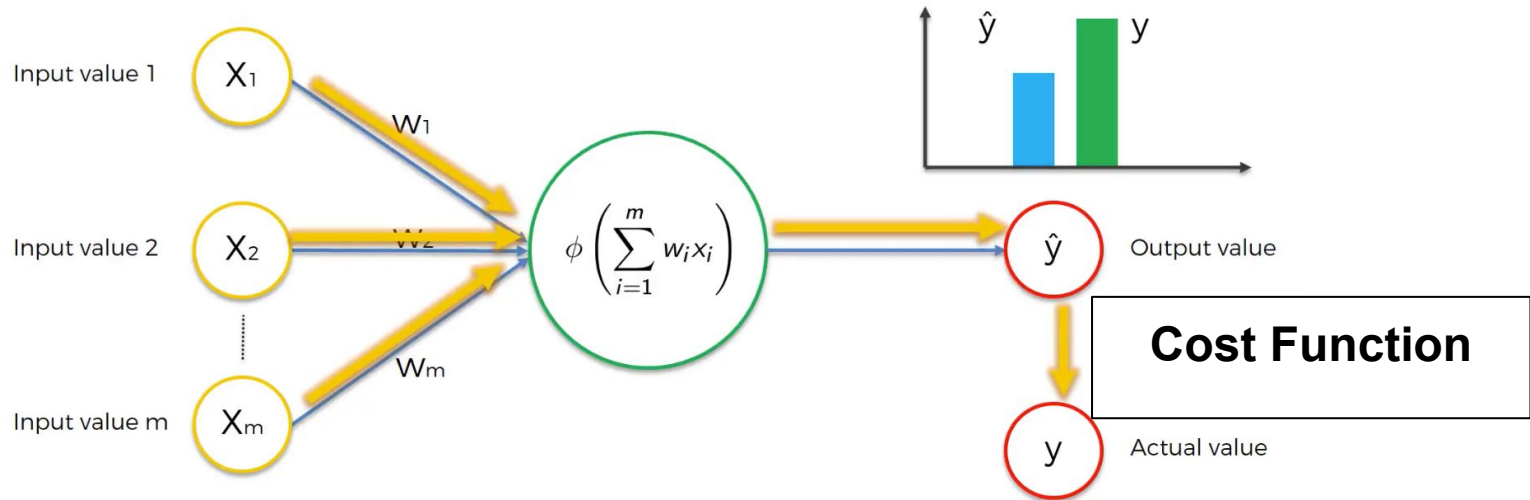


[Reading: The 9 Deep Learning Papers You Need To Know About \(Understanding CNNs Part 3\)](#)

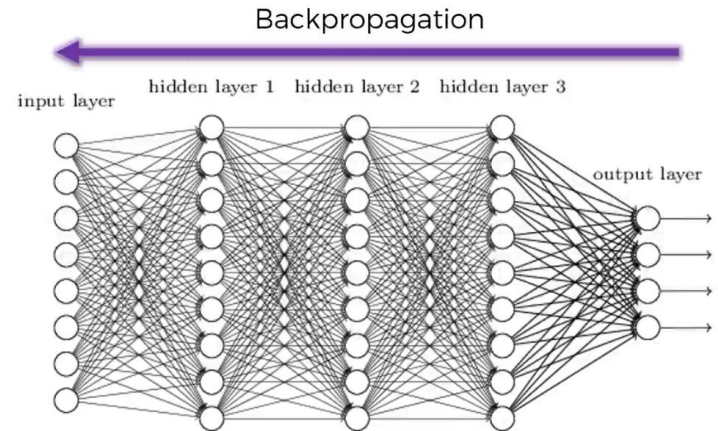
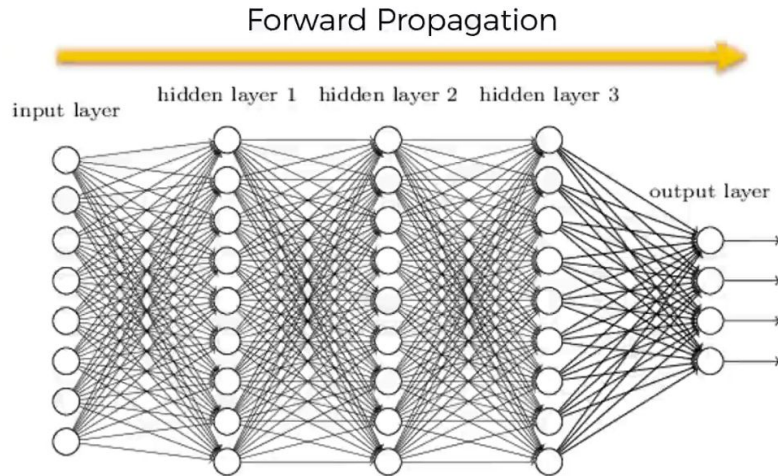
# Fitting the CNN to the images



# How do Neural Networks learn ?

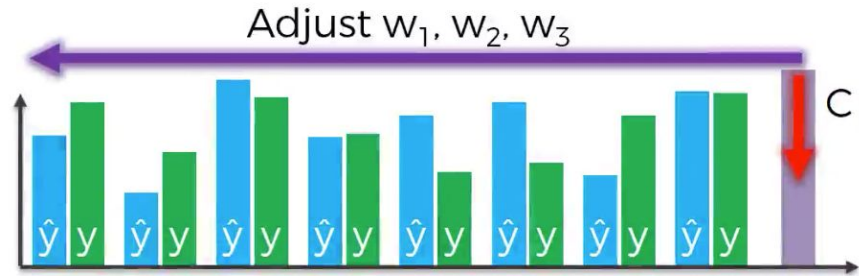
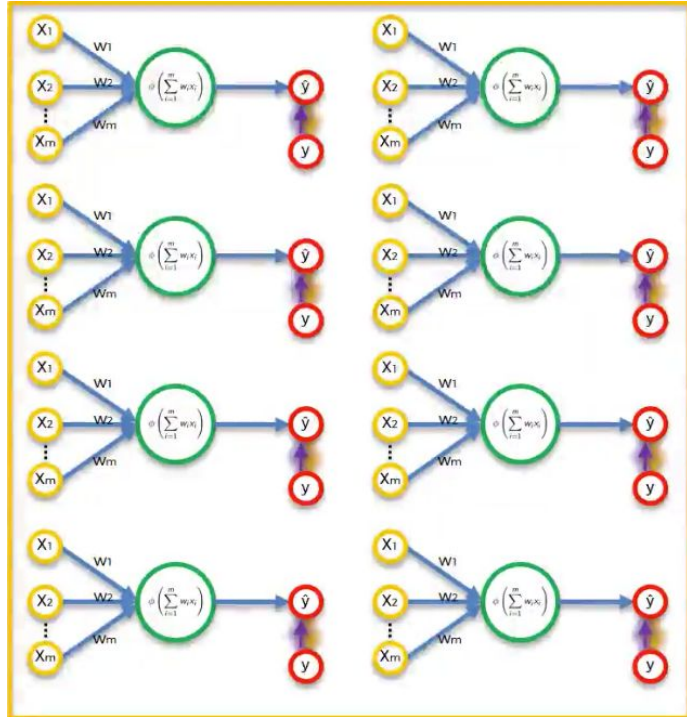


# Backpropagation



[Reading: How the backpropagation algorithm works](#)

# Stochastic Gradient Descent



[Reading: A Neural Network in 13 lines of Python \(Part 2 - Gradient Descent\)](#)

# Training the ANN with Stochastic Gradient Descent

**STEP 1:** Randomly initialise the weights to small numbers close to 0 (but not 0).



**STEP 2:** Input the first observation of your dataset in the input layer, each feature in one input node.



**STEP 3:** Forward-Propagation: from left to right, the neurons are activated in a way that the impact of each neuron's activation is limited by the weights. Propagate the activations until getting the predicted result  $y$ .



**STEP 4:** Compare the predicted result to the actual result. Measure the generated error.



**STEP 5:** Back-Propagation: from right to left, the error is back-propagated. Update the weights according to how much they are responsible for the error. The learning rate decides by how much we update the weights.



**STEP 6:** Repeat Steps 1 to 5 and update the weights after each observation (Reinforcement Learning). Or:  
Repeat Steps 1 to 5 but update the weights only after a batch of observations (Batch Learning).



**STEP 7:** When the whole training set passed through the ANN, that makes an epoch. Redo more epochs.